


Fixed Window




Allow X requests per Y second window.

```
function fixedWindow(maxRequestCount, intervalMs) {  
  let requestCount = 0;  
  let lastTimestamp = 0;  
  
  return () => {  
    const now = Date.now();  
  
    if (now - lastTimestamp > intervalMs) {  
      requestCount = 0;  
      lastTimestamp = now;  
    }  
  
    if (requestCount < maxRequestCount) {  
      requestCount++;  
      return true;  
    } else {  
      return false;  
    }  
  
  };  
}
```

- ✓ Simple to implement
- ✓ Only need to store 2 integers
- ✗ Painful for users when they go over
- ✗ Doesn't allow for bursty traffic
- ✗ Attackers can sustain max throughput

Sliding Window



Allow no more than X requests in any Y second window.

```
function slidingWindow(maxRequestCount, windowMs) {  
  let requests = [];  
  
  return () => {  
    const now = Date.now();  
  
    requests = requests.filter(  
      r => r > (now - windowMs)  
    );  
  
    if (requests.length < maxRequestCount) {  
      requests.push(now);  
      return true;  
    } else {  
      return false;  
    }  
  };  
}
```


✓ Fairer than fixed window

✗ Uses lots of memory

✗ Bursts can exhaust limit for entire window

✗ Attackers can sustain max throughput

Bucket



Bucket overflows if no requests are made for a while

```
function bucketRateLimit(bucketCapacity, intervalMs) {  
  let bucketContents = bucketCapacity;  
  let lastTimestamp = Date.now();  
  
  return () => {  
    const now = Date.now();  
  
    const extraTokens = Math.min(  
      bucketCapacity - bucketContents,  
      Math.floor((now - lastTimestamp) / intervalMs),  
    );  
  
    bucketContents += extraTokens;  
    if (bucketContents === bucketCapacity) {  
      lastTimestamp = now;  
    } else {  
      lastTimestamp += extraTokens * intervalMs;  
    }  
  
    if (bucketContents > 0) {  
      bucketContents--;  
      return true;  
    } else {  
      return false;  
    }  
  
  };  
}
```

- ✓ Only need to store 2 integers
- ✓ Allows for bursts
- ✓ Attackers limited by the interval
- ✗ More complex to explain to API consumers

Exponential Delay

“

Start with an initial delay, then multiply that by a number greater than 1 for each failed passcode attempt.

”

```
function exponential(freeAttemptsCount, initialDelayMs, factor) {
  let attemptsCount = 0;
  let lastTimestamp = Date.now();

  return {
    reset() {
      attemptsCount = 0;
    },

    attempt() {
      const now = Date.now();

      if (attemptsCount < freeAttempts + 1) {
        attemptsCount++;
        lastTimestamp = now;
        return true;
      }

      const delay = initialDelayMs * Math.pow(factor, attemptsCount - (freeAttempts + 1));
      if (now - lastTimestamp > delay) {
        attemptsCount++;
        lastTimestamp = now;
        return true;
      } else {
        return false;
      }
    },
  };
}
```

- ✓ Minimal impact on legitimate users
- ✓ Massive impact on malicious users
- ✓ Only need to store 2 integers
- ✗ Needs some way to reset attempts counter

Recommendations

- ✓ Use Bucket Rate Limiting - for APIs
- ✓ Use Exponential Delay - for password attempts
- ✓ Use the `@authentication/rate-limit` npm package.
- ✓ Combine both where possible
- ✗ Do not use Fixed Window Rate Limiting
- ✗ Do not use Sliding Window Rate Limiting